

In [2]: `import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns`

In [3]: `housing = pd.DataFrame(pd.read_csv("Housing.csv"))
housing.head()`

Out[3]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	13300000	7420	4	2	3	yes	no	no	no	yes	2	yes	furnished
1	12250000	8960	4	4	4	yes	no	no	no	yes	3	no	furnished
2	12250000	9960	3	2	2	yes	no	yes	no	no	2	yes	semi-furnished
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3	yes	furnished
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2	no	furnished

In [4]: `m=len(housing)
m`

Out[4]: 545

In [5]: `housing.shape`

Out[5]: (545, 13)

In [6]: `housing.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  ---
0    price              545 non-null    int64
1    area               545 non-null    int64
2    bedrooms           545 non-null    int64
3    bathrooms           545 non-null    int64
4    stories             545 non-null    int64
5    mainroad           545 non-null    object
6    guestroom          545 non-null    object
7    basement            545 non-null    object
8    hotwaterheating     545 non-null    object
9    airconditioning     545 non-null    object
10   parking             545 non-null    int64
11   prefarea            545 non-null    object
12   furnishingstatus    545 non-null    object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

In [7]: `housing.describe()`

Out[7]:

	price	area	bedrooms	bathrooms	stories	parking
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

In [8]: `varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea']

def binary_map(x):
 return x.map({'yes' : 1, "no" : 0})

housing[varlist] = housing[varlist].apply(binary_map)
housing.head()`

Out[8]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	13300000	7420	4	2	3	1	0	0	0	1	2	1	furnished
1	12250000	8960	4	4	4	1	0	0	0	1	3	0	furnished
2	12250000	9960	3	2	2	1	0	1	0	0	2	1	semi-furnished
3	12215000	7500	4	2	2	1	0	1	0	1	3	1	furnished
4	11410000	7420	4	1	2	1	1	1	0	1	2	0	furnished

In [9]: `from sklearn.model_selection import train_test_split

np.random.seed(0)
df_train, df_test = train_test_split(housing, train_size = 0.7, test_size = 0.3, random_state = np.random)
df_train.shape`

Out[9]: (381, 13)

In [10]: `num_vars = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'price']
df_Newtrain = df_train[num_vars]
df_Newtest = df_test[num_vars]
df_Newtrain.head()`

Out[10]:

	area	bedrooms	bathrooms	stories	parking	price
454	4500	3	1	2	0	3143000
392	3990	3	1	2	0	3500000
231	4320	3	1	1	0	4690000
271	1905	5	1	2	0	4340000
250	3510	3	1	3	0	4515000

In [11]: `df_Newtest.head()`

Out[11]:

	area	bedrooms	bathrooms	stories	parking	price
239	4000	3	1	2	1	4585000
113	9620	3	1	1	2	6083000
325	3460	4	1	2	0	4007500
66	13200	2	1	1	1	6930000
479	3660	4	1	2	0	2940000

In [12]: `def hypothesis(theta, X, n):
 h = np.ones((X.shape[0],1))
 theta = theta.reshape(1,n+1)
 for i in range(0,X.shape[0]):
 h[i] = float(np.matmul((theta, X[i])))
 h = h.reshape(X.shape[0])
 return h`

In [24]: `def gradientDescent(theta, alpha, num_iter, h, X, y, n):
 cost = np.ones(num_iter)
 for i in range(0,num_iter):
 theta[0] = theta[0] - (alpha/X.shape[0]) * sum(h - y)
 for j in range(1,n+1):
 theta[j] = theta[j] - (alpha/X.shape[0]) * sum((h-y) * X.transpose()[j])
 h = hypothesis(theta, X, n)
 cost[i] = (1/X.shape[0]) * 0.5 * sum(np.square(h - y))
 theta = theta.reshape(1,n+1)
 return theta, cost`

In [25]: `def linearRegression(X, y, alpha, num_iter):
 n = X.shape[1]
 column = np.ones((X.shape[0],1))
 X = np.concatenate((column, X), axis = 1)
 theta = np.zeros(n+1)
 h = hypothesis(theta, X, n)
 theta, cost = GD(theta, alpha, num_iter, h, X, y, n)
 return theta, cost`

In [26]: `X_t = df_Newtrain.values[:,[0,1,2,3,4]]
Y_t = df_Newtrain.values[:,5]

X_v = df_Newtest.values[:,[0,1,2,3,4]]
Y_v = df_Newtest.values[:,5]`

In [27]: `mean = np.ones(X_t.shape[1])
std = np.ones(X_t.shape[1])
for i in range(0, X_t.shape[1]):
 mean[i] = np.mean(X_t.transpose()[i])
 std[i] = np.std(X_t.transpose()[i])
 for j in range(0, X_t.shape[0]):
 X_t[j][i] = (X_t[j][i] - mean[i])/std[i]
mean = np.ones(X_v.shape[1])
std = np.ones(X_v.shape[1])
for i in range(0, X_v.shape[1]):
 mean[i] = np.mean(X_v.transpose()[i])
 std[i] = np.std(X_v.transpose()[i])
 for j in range(0, X_v.shape[0]):
 X_v[j][i] = (X_v[j][i] - mean[i])/std[i]`

In [28]: `df_Newtrain.shape`

Out[28]: (381, 11)

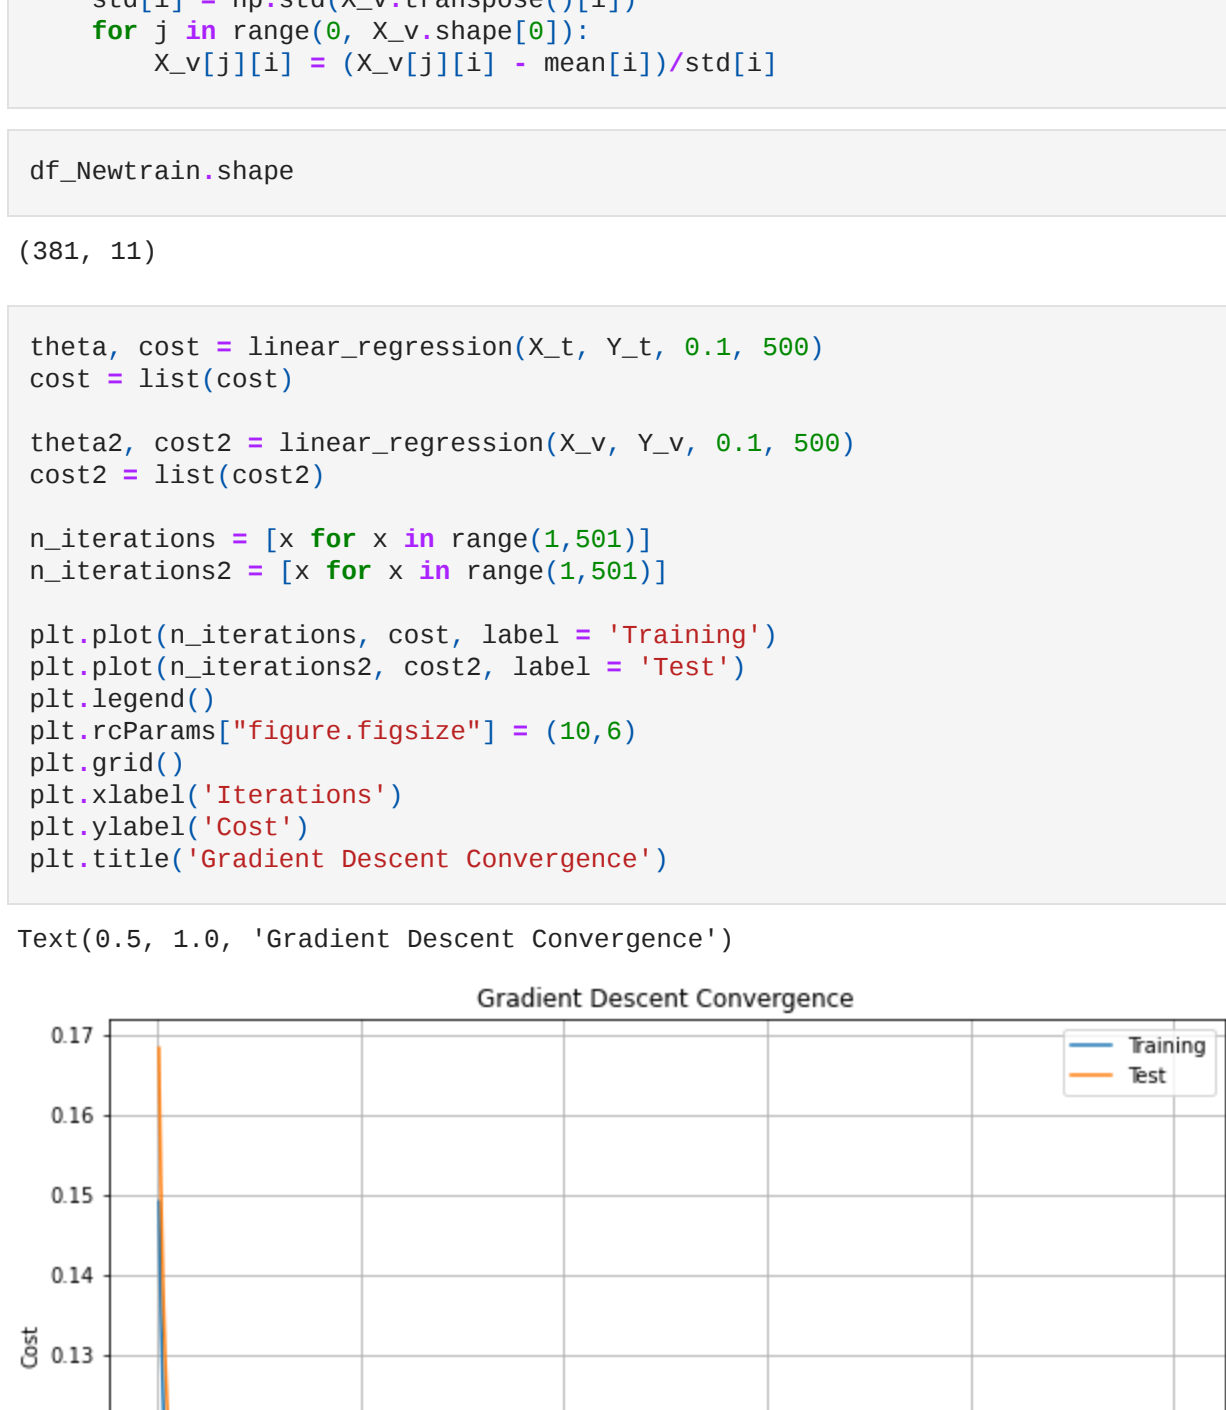
In [29]: `theta, cost = linear_regression(X_t, Y_t, 0.1, 500)
cost = list(cost)

theta2, cost2 = linear_regression(X_v, Y_v, 0.1, 500)
cost2 = list(cost2)

n_iterations = [x for x in range(1,501)]
n_iterations2 = [x for x in range(1,501)]

plt.plot(n_iterations, cost, label = 'Training')
plt.plot(n_iterations2, cost2, label = 'Test')
plt.legend()
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Gradient Descent Convergence')`

Out[29]: Text(0.5, 1.0, 'Gradient Descent Convergence')



In [30]: `num_vars = ['area', 'bedrooms', 'bathrooms', 'mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking', 'prefarea', 'price']
df_Newtrain = df_train[num_vars]
df_Newtest = df_test[num_vars]
df_Newtrain.head()`

Out[30]:

	area	bedrooms	bathrooms	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	price
454	4500	3	1	1	0	0	0	1	0	0	3143000
392	3990	3	1	1	0	0	0	0	0	0	3500000
231	4320	3	1	1	0	0	0	0	0	1	4690000
271	1905	5	1	0	0	1	0	0	0	0	4340000
250	3510	3	1	1	0	0	0	0	0	0	4515000

In [31]: `df_Newtest.head()`

Out[31]:

	area	bedrooms	bathrooms	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	price
239	4000	3	1	1	0	0	0	0	1	0	4585000
113	9620	3	1	1	0	1	0	0	2	1	6083000
325	3460	4	1	1	0	0	0	1	0	0	4007500
66	13200	2	1	1	0	1	1	0	1	0	6930000
479	3660	4	1	0	0	0	0	0	0	0	2940000

In [32]: `X_v = df_Newtest.values[:,0:10]
Y_v = df_Newtest.values[:,10]`

In [33]: `mean = np.ones(X_t.shape[1])
std = np.ones(X_t.shape[1])
for i in range(0, X_t.shape[1]):
 mean[i] = np.mean(X_t.transpose()[i])
 std[i] = np.std(X_t.transpose()[i])
 for j in range(0, X_t.shape[0]):
 X_t[j][i] = (X_t[j][i] - mean[i])/std[i]
mean = np.ones(X_v.shape[1])
std = np.ones(X_v.shape[1])
for i in range(0, X_v.shape[1]):
 mean[i] = np.mean(X_v.transpose()[i])
 std[i] = np.std(X_v.transpose()[i])
 for j in range(0, X_v.shape[0]):
 X_v[j][i] = (X_v[j][i] - mean[i])/std[i]`

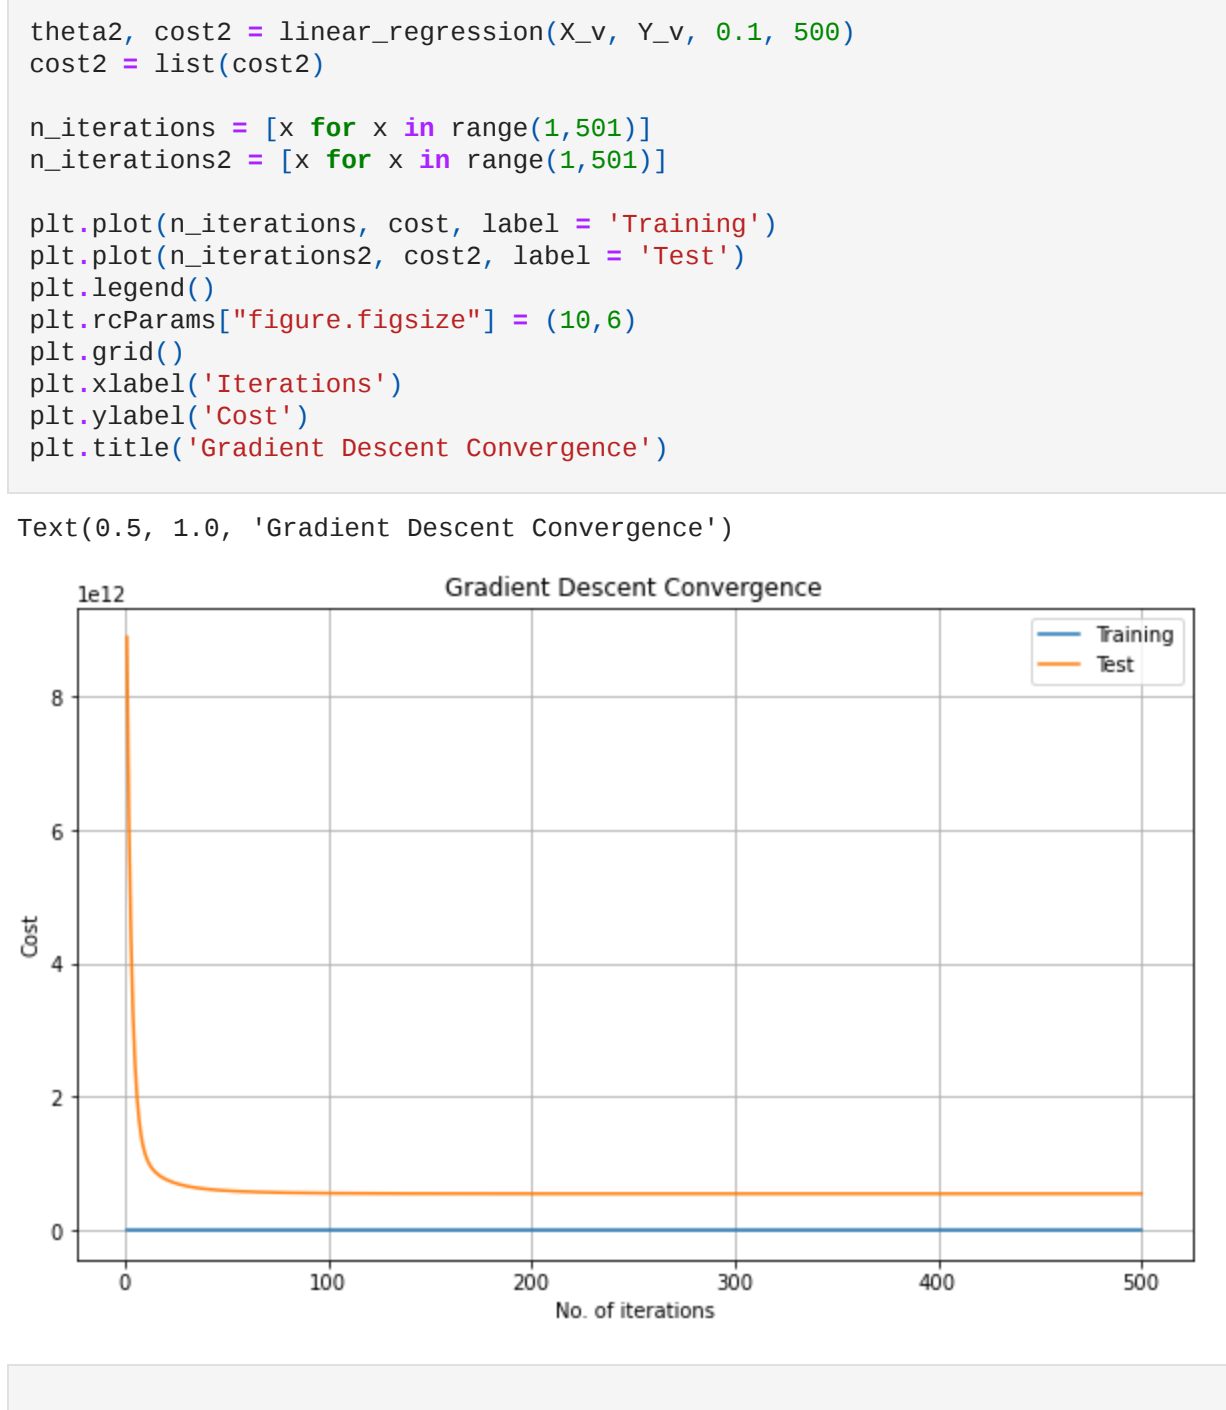
In [34]: `theta, cost = linear_regression(X_t, Y_t, 0.1, 500)
cost = list(cost)

theta2, cost2 = linear_regression(X_v, Y_v, 0.1, 500)
cost2 = list(cost2)

n_iterations = [x for x in range(1,501)]
n_iterations2 = [x for x in range(1,501)]

plt.plot(n_iterations, cost, label = 'Training')
plt.plot(n_iterations2, cost2, label = 'Test')
plt.legend()
plt.rcParams["figure.figsize"] = (10,6)
plt.grid()
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Gradient Descent Convergence')`

Out[34]: Text(0.5, 1.0, 'Gradient Descent Convergence')



In []:

In []: